

AnswerTrust — Speaker Notes

Slides 1 · 2 · 3 | Opening + click-by-click talking points

■ Opening — before Slide 1's first click

"Every enterprise has agentic AI in pilot.
But a fraction of them have it in production.
Not because the models don't work — because **no one can prove the answer**.
That's the gap we close today."

[Pause one beat. Advance to Slide 1.]

Slide 1 — The Gap (3 click stages)

[CLICK 1] — Flow string appears (Pilot → Production → Risk & Audit)

"Here's the journey every agentic AI program is on right now.
You start in **pilot** — exciting demos, real value.
You try to push to **production** — and rollouts pause. Net-new agents freeze.
Then **Risk and Audit** show up with one question: *prove this answer*.
And the room goes quiet."

[Beat. Click.]

[CLICK 2] — Five questions reveal on the left

"These are the five questions Risk and Audit always ask — in this order:
Is the answer **correct**?
Was the data **fit for use**?
Was the user **allowed to see those rows**?
What **exactly** happened — show me the trace?
When do we **intervene** — and at what threshold?

Today, no single Microsoft service answers all five. Purview answers part. Foundry answers part.
Fabric answers part. Defender and Sentinel answer part.
The answer lives in the seams between them — and that seam is empty."

[Beat. Click.]

[CLICK 3] — Hub diagram reveals on the right

"Our proposal is to build the **connective tissue** — the wiring that joins Purview, Foundry, Fabric, and Security into one answerable artifact.
One score per answer. One ledger row. One trace.
We call it **AnswerTrust** — and it's what turns 'we trust the model' into 'here's the proof.'"

Transition → "So how does that work? Four lanes, four jobs."

Slide 2 — Trust Stack (2 click stages)

[CLICK 1] — Four lanes appear (on slide entry)

"Four lanes — every Microsoft service you already own, organized by the job it does for trust.

Observability — App Insights and OTel give us per-hop tracing. Real-Time Hub and KQL give us the AnswerLedger. Eventstream and the Real-Time Dashboard give us live ingest and drill-to-trace.

Governance — Purview labels, DLP, lineage. Entra On-Behalf-Of carries identity end-to-end. Fabric Workspaces and OneLake enforce domain isolation through Azure Policy.

Security — Defender for Cloud watches the agents at runtime. Sentinel correlates abuse. Foundry's red-team agent stress-tests with PyRIT. IRM and Key Vault handle insider risk and secrets.

Quality and Eval — Fabric DQ for validity, completeness, freshness. Foundry evaluators for groundedness, relevance, safety. Golden-set regression on every PR.

Every one of these services already exists. Every one already emits signals. We're not building a platform — we're building the wiring that makes them answer one question together."

[Beat. Click.]

[CLICK 2] — AnswerTrust formula hand + 0.83 chin reveal

"And the wiring produces this:

AnswerTrust = 0.30·Eval + 0.25·DQ + 0.20·LabelCompliance + 0.15·Freshness – 0.10·RedTeamFlags

One number, per answer, on a 0-to-1 scale. Weights are tunable per domain — healthcare may weight DQ higher, finance may weight LabelCompliance higher.

0.83 is the threshold we recommend as a production SLA. Anything below → the answer goes to a human. Anything above → it ships, and the ledger row is the receipt."

Transition → "That single number does something powerful — it gives three different audiences the same artifact."

Slide 3 — Business Value (auto-cascade, no clicks)

Cards animate in

"One artifact. Three audiences.

For the CTO and data leaders — AnswerTrust becomes a publishable SLA. Per agent. Per domain. **Dollars per answer** becomes a steerable metric, not a surprise on the cloud bill. Most importantly — it unblocks the regulated business units that have been sitting on agentic AI for a year because they couldn't get sign-off.

For the data platform team — this is the first end-to-end pattern for observability across Fabric, multi-agent orchestration, and MCP tools. Regression on golden questions. DQ joined to answers via a shared `run_id`. And it reuses the KQL and Eventstream substrate you already operate — no new platform to learn.

For the CISO and compliance — every answer becomes a forensic record. DLP enforcement, DQ verdict, red-team flags, IRM signal — all on one trace. This is *literally* the artifact regulators ask for. DSPM, Sentinel, IRM — one source of truth instead of three dashboards."

Metrics hand

"And here are the SLAs we'll commit to on day one:
100% trace coverage. 100% provenance. 100% DQ coverage. At least 30 golden questions in the regression set. **Under 5 minutes** to detect drift. **Under 10 minutes** to remediate a wrong answer. These aren't aspirational — they're achievable on the existing Microsoft stack the moment we wire it together."

Tanline

"Governance × Quality × Observability × Evaluation × Security — **one ledger row, one number per answer.**"

Transition → "And here's what an operator actually sees when this is live in production..."

■ Delivery cues

Pace: ~110 wpm. Slides 1–3 should run 4–5 minutes total.

Eye contact on the punchlines: "the room goes quiet" · "the seam is empty" · "0.83" · "one ledger row, one number per answer."

Hands: Use the click stages as natural pause points — don't keep talking through the animation.